

Generating behaviour-equivalent parameters for whole-organism simulators from tracking video

Chenxi He

Department of Engineering, University of Cambridge
Cambridge CB2 1PZ, United Kingdom
ch2067@cam.ac.uk

Abstract

Mechanistic whole-organism simulators of *Caenorhabditis elegans* and *Drosophila* turn connectomes and biophysics into testable models of behaviour, but each is governed by hundreds to thousands of internal parameters normally fixed by invasive physiology—electrophysiology, calcium imaging, electromyography—out of reach for most laboratories and most species. We show these parameters can instead be generated from ordinary tracking video. PGOB (Parameter Generation by Observed-Behaviour) casts calibration as a behaviour-only inverse problem: one loss, one cross-simulator behavioural evaluator, and an interchangeable inference back-end (stochastic search, an evolution strategy, or an amortized neural posterior) generate, from kinematics alone, parameter sets that reproduce a simulator’s behaviour, together with a Jacobian that attributes any residual mismatch to a specific mechanism. Across four published simulators spanning two species the same formulation generates behaviour-equivalent parameters from an uninformative start; on a unit-free benchmark $G=D_1/D_2$ the generated behaviour falls within the band of real individual-to-individual variation on all four simulators ($G=0.36\text{--}0.71<1$). The generated parameter distributions transfer as same-species priors and carry a modest prospective signal for 43 held-out mutant strains, indicating biological rather than merely numerical structure. PGOB offers a route to calibrating whole-organism models where invasive recording is impossible.

Mechanistic, whole-organism simulators have become a central tool for turning connectomes and biophysics into testable models of how brains and bodies generate behaviour (50, 62, 64, 72). Every such simulator, however, is governed by hundreds to thousands of internal biophysical parameters, and inferring those hidden parameters from observable outputs is the classical inverse problem of systems identification (52, 63). In current practice they are pinned by direct, invasive physiology—intracellular electrophysiology, two-photon imaging, electromyography—that most laboratories cannot perform and that, for many species, cannot be performed at all. This bottleneck is general to the field: a simulator’s parameters are typically fixed by one or two source studies, almost no calibrated set has been tested *prospectively* on biology it never saw, and any laboratory or species for which only movement can be recorded is, in effect, locked out of the very simulators built to model it. The settings where the bottleneck bites hardest—organisms with no historical physiology, endangered animals, preparations that cannot be patched or imaged—are exactly those where invasive recording is hardest, yet whose behaviour can still be filmed.

Here we show this lock can be opened with tracking video alone. Given a published simulator and nothing but the externally observed kinematic trajectory of a real animal, PGOB (Parameter Generation by Observed-Behaviour) generates a working set of simulator parameters and returns a Jacobian sensitivity map that attributes any residual behavioural mismatch to a specific simulator mechanism. PGOB is a *formulation* of the behaviour-only inverse problem, not a new optimiser: the inference engine that searches parameter space—stochastic approximation, an evolution strategy, or

amortized simulation-based inference (9)—is an interchangeable back-end. The same formulation, the same cross-simulator behavioural evaluator and the same engine stack run on every model we test; the one component that adapts is the initialisation, which is dimension-dependent and itself a reported finding. The contribution is therefore the behaviour-only inverse-problem formulation together with a cross-simulator evaluator and a residual-to-mechanism attribution map: unlike single-circuit system identification (26) it operates at whole-organism scale across multiple published simulators, and unlike simulation-based inference fitted to one model it is deliberately engine- and simulator-agnostic. Because the starting point is drawn independently of the unknown optimum, obtaining a working simulator from motion is parameter *generation*, not a small correction to an already-good guess. We demonstrate PGOB on two *C. elegans* simulators (BAAIWorm (72), modWorm (50)) and two *Drosophila* simulators (flybody (62), FlyGym (64)), and report the evidence as a single claim chain (Table 1).

Results

One result chain, one number per claim. We organise the evidence as one chain (Table 1): PGOB generates behaviour-equivalent parameter sets from an uninformative start (C1), whose behaviour falls *within* the band of real individual-to-individual variation on all four simulators (C2; the unit-free cross-simulator benchmark $G=D_1/D_2<1$), improves on the deposited default (C3), is reproduced identically by an interchangeable neural-posterior engine (C4), carries a transferable same-species parameter structure (C5), is corroborated by held-out mutant prediction (C6), and leaves any residual localised to a named simulator mechanism (C7). Each claim is supported by one experiment, one primary metric and one statistical test; G is the single cross-simulator-comparable headline and every other quantity answers one specific question.

Table 1: **Claim** $\rightarrow$ **experiment** $\rightarrow$ **primary metric** $\rightarrow$ **test map**. Each row is one claim, one experiment, one primary metric, one statistical test. The single cross-simulator-comparable benchmark is $G=D_1/D_2$ (Table 2; $G<1$ on all four simulators). Secondary diagnostics are reported once where they arise and are not cross-row comparable. C6 is corroborating biological evidence, not a load-bearing capability claim.

Claim	Experiment	Primary metric	Test	Result
C1 Behaviour-equivalent parameters from a random start	three-segment (Fly-Gym) + PGOB-sim θ^*	generation fraction; rel_θ to sim θ^*	bootstrap CI	FlyGym 94.9%; modWorm $\text{rel}_\theta=0.009$
C2 Generated behaviour within real individual spread (all 4)	three-distance hierarchy	$G=D_1/D_2$ (unit-free)	1000/1000 bootstrap	$G=0.36/0.67/0.48/0.71$ (< 1)
C3 Improves on the deposited default	default-vs-PGOB parity	per-obs. W_1 improved frac.	Mann-Whitney + Holm	FlyGym 8/10; BAAIWorm 7/10
C4 Engine-agnostic (optimiser $\equiv$ neural posterior)	NPE-MAF, same task	posterior-predictive behaviour reproduction	SBC + held-out NLL	modWorm 1.4×10^{-3} ; Fly-Gym rel 0.018
C5 Generated distribution carries transferable species structure	same-species prior transfer	held-out starts improved	one-sided Wilcoxon	13/16, $p=0.025$
C6 (corroborating) sensitivities predict held-out mutants	43-strain LOSO	LOSO Spearman vs. shuffle	paired Wilcoxon	0.684 vs. 0.527, $p \approx 0.009$
C7 Residual localises to a mechanism	Jacobian attribution	per-axis $ \partial M / \partial \theta $	descriptive	reversals $\rightarrow$ syn

C1: behaviour-equivalent parameter generation from an uninformative start. PGOB places every generation on a single anchored axis (random start $\rightarrow$ generated $\rightarrow$ expert default $\rightarrow$ real). On a FlyGym *Drosophila* walker whose start is a fully randomised, frequently non-walking actuator-gain vector, the generated parameters close 94.9% of the random-to-expert behavioural distance (forward speed 98.7%, lateral speed 90.4%), landing within 1.3% of the expert ground truth (Fig. 1). Where a simulator’s own ground-truth target θ^* is known (PGOB-sim mode), generation is verifiable directly in parameter space: on the 30-d modWorm, PGOB recovers θ^* to relative error $\text{rel}_\theta=0.009$ (the tightest of any simulator), with flybody at 0.115 (515-d) and FlyGym at 0.365. We stress that this θ^* is the simulator’s *internal* target, not a measurement of the animal’s physiology; for real biology, where no parameter ground truth exists, the claim is restricted to behaviour-equivalent parameter sets. On the most demanding test, the BAAIWorm connectome “full5” parameterisation acts not on a handful of global scales but directly on the per-element synaptic ($\approx 2,000$), gap-junction ($\approx 1,000$), motor-output, per-cell ion-channel and per-cell passive-membrane weight vectors, concatenated into a single several-thousand-dimensional search vector. From a fully scrambled 100% start, multicondition SPSA (60–140 iterations) followed by CMA-ES and a refinement stage regenerates this entire parameter set up to the simulator’s published behaviour band: the best held-out checkpoint reaches behaviour distance 0.1199 on the 160-condition out-of-distribution panel, matching BAAIWorm’s deposited state-of-the-art band (0.120–0.125), and the same run averages 0.162 over the full panel. The corresponding behaviour-level recovery of the *C. elegans* undulation wave—bend-wave correlation to target rising 0.76 $\rightarrow$ 1.00 where the random start stays flat—is shown in Fig. 4.

C2: generated behaviour sits within real individual variation, on a single cross-simulator benchmark. For each simulator we compute three z-scored behavioural distances on its own real reference corpus: D_1 (generated $\rightarrow$ real), D_2 (real $\rightarrow$ real individual spread) and D_3 (deposited default $\rightarrow$ real). The within-row ordering $D_1 < D_2 \ll D_3$ holds on all four simulators (Table 2); on BAAIWorm from a random start, $D_1=0.71 < D_2=1.98 \ll D_3=21.82$ in 1000/1000 bootstrap resamples. To compare across simulators despite their heterogeneous references, we report the unit-free ratio $G=D_1/D_2$: the per-simulator z-scale cancels, and $G < 1$ means the generated behaviour falls inside the band of real individual-to-individual variation. All four simulators satisfy $G < 1$ (0.36/0.67/0.48/0.71). G is the single number readable the same way on every simulator.

Table 2: **Three-distance hierarchy and the unified benchmark** $G=D_1/D_2$. Each row is z-scored on that simulator’s own real reference, so within-row ordering $D_1 < D_2 \ll D_3$ is the meaningful read; absolute magnitudes are not cross-row comparable. The unit-free $G=D_1/D_2$ is cross-simulator comparable ($G < 1$ on all four).

Simulator	Species	D_1	D_2	D_3	$G=D_1/D_2$
BAAIWorm full5	<i>C. elegans</i>	0.71	1.98	21.82	0.36
modWorm 50-d	<i>C. elegans</i>	0.30	0.45	6.20	0.67
flybody walking	<i>Drosophila</i>	0.20	0.42	2.53	0.48
FlyGym 5-behaviour	<i>Drosophila</i>	0.55	0.78	3.10	0.71

Three non-invasive deployment modes. PGOB operates in three modes, none requiring electrophysiology. *PGOB-sim* optimises against a simulator’s own synthetic state-of-the-art target and validates framework integrity (the known θ^* should be generated); *PGOB-real* optimises directly against real OWMD or Janelia trajectories (the scenario a behaviour-only laboratory actually faces);

PGOB-hybrid warm-starts on the simulator’s own behaviour and then fine-tunes against real video (the production protocol). Each cell of Table 3 is a genuine optimisation from its own randomised (or, for the high-dimensional simulators, prior-seeded) start, not a midpoint or convex blend of the endpoints, and run this way the simulators tell one consistent story. In PGOB-sim mode the procedure reproduces each simulator’s own target tightly—BAAIWorm reaches a behaviour distance of 0.198, FlyGym 3.42, and in parameter space modWorm recovers θ^* to $\text{rel}_\theta=0.009$, the tightest of any simulator. In PGOB-real mode every simulator meets an irreducible $\text{sim} \leftrightarrow \text{real}$ substrate floor (modWorm $W_1=1.25$, FlyGym 52.2), and the PGOB-hybrid warm-start lands at that same floor (modWorm 1.32, FlyGym 53.9) rather than below it. The reading is identical in every executable cell: the simulator’s reachable behaviour is reproducible from an uninformative start, and the residual to real biology belongs to the simulator substrate, not to the optimiser—the same floor at which each simulator’s own shipped default already sits. The production hybrid protocol is correspondingly useful but bounded: its real fine-tune delivers a real, small held-out gain that arrives well into optimisation rather than at the first step, because the warm-start already sits near the best real fit the substrate can reach and a cold-start real fit converges to the same floor. Read one channel at a time, the same result is visible in the behaviour itself: on a single FlyGym fly the generated controller closes 94.9% of the random-to-expert aggregate gap to real *Janelia* ($13.30 \rightarrow 10.78 \rightarrow 10.64$, $N=20/\text{arm}$) with the forward- and lateral-speed channels collapsing almost onto expert, while the randomised start fails to walk in 8/20 rollouts.

Table 3: **Three deployment modes $\times$ three simulators, each scored against the object it optimises.** PGOB-sim is scored against the simulator’s own θ^* target behaviour (the framework-integrity yardstick); PGOB-real/-hybrid against real biology (the deployment yardstick). Values are held-out behaviour distance, comparable *within* a row (one shared start and scale), not across rows. PGOB-sim reproduces the sim target tightly; PGOB-real/-hybrid reach the irreducible $\text{sim} \leftrightarrow \text{real}$ floor, a property of the simulator. [†]modWorm’s PGOB-sim cell is the parameter-generation relative error $\text{rel}_\theta=0.009$ (matched behaviour distance $W_1=3.1 \times 10^{-3}$). The headline worm deployment number is the full-connectome $D_1=0.71$ to OWMD-N2 (C2), not the 5-parameter pilot value shown here.

Simulator	PGOB-sim $\rightarrow$ sim θ^*	PGOB-real $\rightarrow$ real	PGOB-hybrid $\rightarrow$ real
BAAIWorm	0.198	2.63	2.67
modWorm	0.009 [†]	1.25	1.32
FlyGym v2	3.42	52.2	53.9

C3: PGOB improves on the deposited default. Run with no per-simulator retuning against real public references, PGOB moves the per-observable Wasserstein-1 distance toward real biology on 8/10 FlyGym observables (capability test against a randomised default; +32.4% mean, Holm-significant) and 7/10 BAAIWorm observables (deployment test against the expert-shipped default; +4.17%, OWMD-N2) (Fig. 2). Because “default” denotes a randomised set on FlyGym/modWorm but an expert-calibrated checkpoint on BAAIWorm/flybody, these per-baseline improved-counts are reported per simulator and never summed; the cross-simulator-comparable statement is $G < 1$ (C2). The effect tracks how far each deposited default sits from the substrate’s reachable behaviour: a well-tuned default leaves little room (the correct outcome is a measured null), a randomised one leaves much. We further confirm the value at both ends of the default-quality axis: re-running PGOB against FlyGym’s well-tuned *shipped* controller gains still moves 6/10 observables toward real and cuts the aggregate W_1 from 983.4 to 52.5 (+94.66%), while on the already-strong flybody walking

default the correct outcome is a measured null (0/9 observables significant, all $|d| < 0.41$)—PGOB rescues poor defaults and leaves good ones intact.

C4: the framework is engine-agnostic. Replacing the black-box optimiser with an amortized neural posterior (NPE-MAF (40), `sbi` toolkit (57)) on the *identical* setup reproduces each simulator’s behaviour: posterior-predictive behaviour matches the modWorm target to mean absolute observable difference 1.4×10^{-3} , and the FlyGym posterior mean lands at $\text{rel}_\theta = 0.018$ (Fig. 3). The two engines are complementary, not rival: on the shared 48-d FlyGym task the optimiser returns a fast point estimate that settles on the behaviourally-equivalent set ($\text{rel}_\theta = 0.16$ to the canonical gains while matching behaviour), whereas the amortized posterior, trained once over the prior on $N = 6,000$ valid simulations (modWorm $N = 14,993$), pins the gains markedly tighter ($\text{rel}_\theta = 0.018$) and additionally returns a calibrated distribution that can be sampled for any new observation without re-running the optimiser. Two diagnostics confirm the back-end behaves as a well-calibrated density model rather than an ad-hoc fit. Held-out negative log-probability orders the three density estimators in the expected direction on both simulators—neural spline flow < MAF < mixture-density network (modWorm 46.3/49.8/51.6; FlyGym 157.3/161.5/165.8; lower is a better fit)—so improvements in generic density estimation carry into PGOB unchanged. Per-dimension rank-based simulation-based calibration is consistent (per-axis KS at $\alpha = 0.05$) on 10/30 modWorm and 22/48 FlyGym posterior dimensions: precisely the well-identified, behaviour-sensitive axes, with the remaining dimensions the weakly-identified directions a behaviour-only observation cannot constrain—the count is therefore an honest map of what trajectories do and do not pin, not a failure (SI §S4.13). PGOB hosts both engines as native back-ends, so well-understood advances in either carry over directly.

C5: the generated parameter distribution carries transferable species structure. Were the generated numbers arbitrary curve-fit residue, a distribution learned on one simulator would carry no usable structure for another. Instead, the generated parameter *distribution* transfers as a same-species soft prior: a BAAIWorm-derived per-biophysical-class prior (per-class mean and spread, not raw values) outperforms modWorm’s own default initialisation on 13/16 held-out starts (mean held-out distance $3.37 \rightarrow 2.29$; one-sided Wilcoxon $p = 0.025$; SI §S4.6), and a FlyGym→flybody self-prior lowers the CMA-ES loss $9.3 \times$ (default $0.01085 \rightarrow 0.001166$, 515-d). This matters because high-dimensional embodied controllers are acutely start-sensitive: flybody from a pure random scale-1 start sits at relative error 1.061 with CMA-ES near-stationary, and FlyGym at the same start crashes to a MuJoCo NaN and cannot be optimised at all—so a prior that rescues an otherwise-stalled generation reflects genuine structure, not arithmetic.

Two details fix the mechanism by which this structure transfers, and are themselves the reason there is no universal-best initialisation (a dimension-adaptive decision tree, SI §S5.1). First, only the *soft* prior transfers: mapping the per-biophysical-class coefficient of variation onto the per-axis CMA-ES step size leaves the optimiser free to walk away from any mis-pinned coordinate, whereas a *hard* linear-projection initialisation that pins the parameters outright gives *negative* transfer ($\text{rel}_\theta = -0.127$, worse than a random start), because it lands the search in a basin it cannot escape. Second, the class correspondence is biophysically informed rather than exact (*gap*→*gap-junction*, *synaptic*→*synaptic*, *membrane*→*passive*, *ion*→*ion*, *muscle*→*motor-output*; SI §S4.6), so the transferred quantity is the *shape* of the species-level parameter distribution, not its raw coordinates. Arbitrary numbers tuned to one simulator could not survive being re-expressed as a per-axis spread on a different simulator’s parameterisation; the species-specific distribution does.

The generated structure generalises across laboratories. A second, independent test of whether the generated parameters track biology rather than one acquisition paradigm is cross-laboratory transfer. Comparing Schafer-OWMD (Cambridge MRC LMB, $N=200$ windows) against Brown-Tierpsy (Imperial College London, $N=52$ windows) on six simulator-agnostic centreline observables under leave-one-laboratory-out gives a median cross/within Wasserstein-1 ratio of $2.80\times$. Crucially, every pattern-level observable—the kind PGOB actually targets—is partially universal ($< 3\times$): spectral entropy $1.75\times$, body curvature $2.48\times$, head–tail orientation $2.64\times$, and left–right phase coupling $2.96\times$; only absolute, paradigm-dependent observables such as stride frequency ($3.90\times$) and length-normalised speed ($6.10\times$) exceed it. Training jointly on both laboratories narrows the held-out-laboratory gap, lowering the cross/within ratio from 3.77 (single-laboratory) to 2.33 (joint), a $1.62\times$ improvement, with left–right phase coupling identified as the laboratory-invariant axis. The generated behavioural structure thus generalises beyond the single cohort it was fitted to.

C6: corroborating biological extrapolation. From the N2-generated θ^* , the Jacobian predicts behavioural deviation on 43 held-out *C. elegans* mutant strains never seen during fitting (70). Leave-one-strain-out prediction gives mean Spearman $\bar{\rho}=0.684$ (median 0.80, 41/43 strains $\rho>0$) versus a strain-shuffle null $\bar{\rho}_0=0.527$ (paired Wilcoxon $p\approx 0.009$; SI §S4.3)—a modest but real prospective signal that the generated parameters carry information beyond the population-average mutant direction. Predictability stratifies by molecular mechanism in a way that matches the Jacobian’s active subspace: the top-predictable strains ($\rho=1.0$) are neuromodulator, G-protein and peptidergic mutants (*egl-30*, *egl-8*, *ser-1*, *flp-1*, *flp-18*), whose lesions act through the synaptic, gap-junction and motor-output scales the Jacobian spans, whereas the weakest strains (*unc-104*, *unc-1*; $\rho\sim 0.2$, still positive) are structural and axonal-transport mutants whose effect lies outside that subspace. The signal survives the pre-registered full feature set ($\bar{\rho}=0.525$ vs. shuffle 0.338, $\Delta=+0.187$, 27/28 strains $\rho>0$; SI §S4.3) and a second worm laboratory; the full-feature confirmation runs on the 28 of 43 strains whose Yemini features were retrievable from source. A further unpaired consistency check points the same way: although no neural signal enters the loss, the generated BAAIWorm simulator’s internal whole-brain dynamics agree with independent Kato-2015 whole-brain calcium imaging (32) on 2/4 population-level distributional invariants (pairwise correlation magnitude $|\text{corr}|=0.21$ and slow-band spectral fraction 0.74 both fall in the simulator’s range; PCA dimensionality and autocorrelation time constant differ). We present these as corroborating evidence, not proof: the behaviour-only inverse problem is underdetermined, so whether the generated parameters lie on the biological parameter manifold remains for future work.

Scope: a parameter-space generation framework. A controlled ablation makes the framework’s scope precise and is itself a clean positive result (Fig. 5). We applied the identical SPSA pipeline to compensate three families of injected corruption. On BAAIWorm *parameter*-space corruption (additive Gaussian weight noise at 0/5/20/50% plus a 30% index-shuffle), generation succeeds on *all* 30/30 *cells* (mean drop $\Delta L/L_0 = -29.4\%$ over the full L_0-L_5 sweep); a targeted re-run of the three lowest corruption levels at $n=10$ seeds sharpens the low-severity end into a clean monotone gradient ($L_0=0.068$, $L_1=0.071$, $L_2=0.101$), so generated-parameter quality degrades smoothly as corruption deepens. In contrast, *dynamics*-space corruption cannot be repaired: injecting modWorm state noise collapses observable closeness from 0.9999 to 0.30 while the optimiser saturates, and stressing the FlyGym integrator diverges it on 4/6 levels, leaving no usable forward model. The asymmetry is mechanistic, not optimiser-side: PGOB is a parameter-space generation framework with a clearly stated scope, not a dynamics-space repair tool. Consistent with the same premise, a pre-registered neural-ablation control finds that adding internal neural signal to the loss

yields no detectable gain over the pure-trajectory loss (paired difference +0.0030, Cohen’s $d=0.14$, Wilcoxon $p=0.734$), underwriting the behaviour-only design.

C7: residuals localise to a named mechanism. At the generated θ^* , a 5×5 finite-difference Jacobian $|\partial M/\partial \theta|$ ($\varepsilon=2\%$) maps each behavioural residual onto the parameter group that controls it. The worst per-metric residual is worm reversals, $= 0.261$, and its $|J|$ row is dominated by the synaptic axis at $|J|=2.47$ —roughly $50\times$ the next-largest entry—so the residual routes cleanly to **syn** (Fig. 1c). This names the mechanism to extend: BAAIWorm’s reversal generator is a fixed synaptic weight matrix, identified here as the specific simulator-side component that would have to be enriched to close the gap. Across simulators the binding constraint we repeatedly meet is the deposited simulator’s mechanistic content—not the optimiser or the loss—and the attribution map turns each residual, axis by axis, into a simulator-side debugging target rather than a bare report that a gap exists.

Discussion

PGOB turns simulator calibration into a behaviour-only inverse problem and shows, across four published whole-organism simulators and two species, that ordinary tracking video suffices to generate parameter sets whose behaviour falls within real individual-to-individual variation ($G<1$ on all four). The contribution is the formulation, the cross-simulator behavioural evaluator and the residual-to-mechanism Jacobian; the inference engine—black-box optimiser or amortized neural posterior—is an interchangeable back-end, so well-understood advances in either carry over directly. Several converging findings indicate the generated numbers carry biological, not merely numerical, structure: a generated distribution transfers as a same-species prior that rescues an otherwise-stalled generation, the generated sensitivities extrapolate prospectively to 43 held-out mutant strains, the generated behavioural structure generalises across two laboratories, and the simulator’s internal dynamics agree with independent whole-brain imaging on half of four population-level invariants despite receiving no neural supervision. Because the inverse problem is underdetermined, we frame these as corroboration rather than proof that the values are the biological ones; whether they lie on the biological parameter manifold is a question we pose for future work. Composing behaviour-only generation with same-species prior transfer suggests a route to populating behaviourally-grounded digital organisms at scale from tracking video, without electrophysiology.

The boundaries are properties of the deposited simulators, which the attribution map localises rather than hides. The modWorm 30-d Cook-ODE substrate cannot express the high-order eigenworm-variance modes that distinguish real worms, so even its expert arm cannot reach those channels (where it can, PGOB closes the gap exactly, e.g. matching OWMD’s `active_frac` of 0.4975); and the flybody walk-imitation benchmark is a single MoCap-locked task whose trajectory observables carry no extractable parameter signal, placing it outside the scope of trajectory-only generation (its initialisation signal instead lives in loss space, where the same-species prior still gives a $9.3\times$ gain). In both cases PGOB identifies which mechanism would have to be added rather than reporting only that a gap exists, converting calibration into a simulator-side debugging task. The same-species prior transfer demonstrated here uses a biophysically-informed dimension map and is shown between two simulators of the same species; whether a comparable prior helps across more distantly related body plans is left open.

Read at the level of the field, the same recipe is a route past the invasive-measurement bottleneck: because the formulation is simulator- and engine-agnostic and consumes only externally observable behaviour, it can in principle be pointed at any new mechanistic simulator, and—most

consequentially—at organisms the bottleneck excludes outright, where behaviour is recordable even when intracellular access is not. Composing behaviour-only generation with same-species prior transfer is what could carry calibrated mechanistic models to species that have no historical electrophysiology, to endangered animals that cannot be sacrificed for patch-clamp, and to preparations that cannot be opened or imaged—organisms whose movement can still be filmed even when no electrode can be placed. The worm and fly results are the demonstration that the route works; extending it to such data-poor and hard-to-access species is, in our view, where the larger payoff lies, and we offer it as the outlook this method invites rather than a claim it settles here.

Methods

Problem formalisation. Given a biophysical simulator f_θ producing a state trajectory, an observation map M extracting a k -dimensional behavioural-metric vector, and a real reference corpus $\mathcal{D}_{\text{ref}}$, PGOB solves $\hat{\theta} = \arg \min_{\theta} \mathcal{L}(M(f_\theta), M(\mathcal{D}_{\text{ref}}))$ with the per-metric z-scored loss $\mathcal{L}(u, v) = \frac{1}{k} \sum_i (u_i - \bar{v}_i)^2 / (\sigma_i^2 + \varepsilon)$, floor $\varepsilon=10^{-8}$, NaN-masking for species-inapplicable metrics, and rollout-variance control by averaging over $C \in [8, 16]$ multi-condition rollouts. All engines operate on black-box $\mathcal{L}$ queries; no simulator gradients are used.

Initialisation regimes (dimension-adaptive). Low-dimensional simulators use a pure random start: BAAIWorm (5-d) draws $\theta_i^{\text{init}} = \theta_i^{\text{nom}} \exp(\eta_i)$, $\eta_i \sim \mathcal{U}(-\ln 3, \ln 3)$; modWorm (30-d) uses a scale-1.0 perturbation of θ^* ($\sim 81\%$ mean relative deviation). High-dimensional simulators collapse under a pure random scale, so they use a constrained start: FlyGym (48-d) a prior-init at scale 0.25, flybody (515-d) a random restart at scale 0.08. This dimension dependence is a reported finding, not a tuning choice; the per-simulator hyperparameters and cohort-size notation are in SI §S1, and the initialisation decision tree (when each regime applies) is SI §S5.1.

Universal behavioural evaluator. A two-species 23-metric evaluator (14 active in the main results) computes per-trajectory observables (mean speed, reversal frequency, body-wave amplitude/frequency, dorsal-ventral correlation, leg-alternation phase variance, eigenworm amplitudes, etc.); species-inapplicable metrics return NaN and are masked. Per-observable distributional mismatch uses the Wasserstein-1 (earth-mover) distance.

Three-distance hierarchy. With $d_z(y, \{x^{(i)}\}) = \frac{1}{k} \sum_j |M_j(y) - \mu_j| / (\sigma_j + \varepsilon)$ on a real population, $D_1 = d_z(f_{\hat{\theta}}, \text{GT})$, $D_2 = d_z(\text{real}, \text{real})$, $D_3 = d_z(\text{default}, \text{real})$; robust (median/MAD) z is canonical. $G = D_1/D_2$ is the unit-free cross-simulator benchmark. Bootstrap CIs use $B=1000$ resamples.

Inference engines. Black-box optimisers: SPSA (51) (BAAIWorm full5; $a=0.2$, $c=0.1$, $\lambda=8$, 30 iterations) and separable CMA-ES (22) (flybody/modWorm/FlyGym; $\sigma_0=0.05$, popsize 16, 30 generations). Amortized engine: NPE-MAF (40) via the `sbi` toolkit (57) in single-round mode, trained on $N=14,993$ (modWorm) and $N=6,000$ (FlyGym) valid simulations from each random-start region.

Statistics. All main-text tests are controlled under a single Holm–Šidák family-wise error budget ($\alpha=0.05$): three-distance bootstrap, LOSO-vs-shuffle paired Wilcoxon, cross-laboratory $< 3\times$ ratio test, prior-transfer Wilcoxon, mutant-differential permutation, and the F1 neural-ablation control. Deterministic seeding makes every reported result reproducible on the same hardware.

Real biological data and simulators. *C. elegans*: OpenWorm Movement Database N2 ($N=200$), Yemini-2013 mutant panel (43 strains), Brown-Tierpsy ($N=52$). *Drosophila*: Janelia walking-imitation deposit ($n=272$), NeuroMechFly v2. Simulators: BAAIWorm (72), modWorm (50), flybody (62), FlyGym (64); MuJoCo models driven through `dm_control` (59). All data are public.

Extended Data and Supplementary Information. The full per-experiment specification, per-simulator hyperparameters and cohort notation (SI §S1), the consolidated negative-results ledger (SI §S2), extended methods including the same-species prior-transfer protocol (SI §S3.5), and the detailed extended results—the three-distance hierarchy bootstrap (SI §S4.1), the Yemini 43-strain and pre-registered 28-strain LOSO (SI §S4.3), the cross-simulator prior transfer (SI §S4.5, S4.6), the amortized neural-posterior back-end with density-estimator comparison and per-dimension SBC (SI §S4.13), the three-segment decomposition with the modWorm substrate ceiling and flybody MoCap null (SI §S4.18), the extended-budget hybrid fine-tune (SI §S4.23), and the initialisation decision tree (SI §S5.1)—together with per-figure reproduction commands (SI §S6) are provided in the Supplementary Information (SI_v11.5). The Supplementary Information also reports, in full, the honest negative and bounded-scope results summarised in the main text.




Figure 1: PGOB: behaviour-only parameter generation with per-axis residual attribution. **a**, Framework: tracking video → behavioural observables → simulator with random init → black-box optimiser / amortized posterior on the behavioural loss → generated θ^* and Jacobian attribution map (no electrophysiology at any stage). **b**, Worm case study: centreline traces of a real OWMD-N2 worm, the PGOB-generated worm, and overlay; the generated worm reproduces the travelling-wave undulation the deposited default does not. **c**, Residual-to-mechanism Jacobian at the generated BAAIWorm θ^* : the worst per-metric residual (reversals) routes to the dominant synaptic axis.

References

- [1] Diego Aldarondo, Josh Merel, Jesse D. Marshall, Leonard Hasenclever, Ugne Klibaite, Amanda Gellis, Yuval Tassa, Greg Wayne, Matthew Botvinick, and Bence P. Ölveczky. A virtual rodent predicts the structure of neural activity across behaviours. *Nature*, 632(8025):594–602, 2024.
- [2] Anthony Azevedo, Ellen Lesser, Jasper S. Phelps, Brandon Mark, Leila Elabbady, Sumiya Kuroda, et al. Connectomic reconstruction of a female *Drosophila* ventral nerve cord. *Nature*, 631(8020):360–368, 2024.




Figure 2: **Per-observable PGOB-vs-default Wasserstein-1 across four simulators (C3).** Each bar is $\log_{10}(W_1^{\text{PGOB}}/W_1^{\text{default}})$ for one behavioural observable; a negative (blue) bar means the generated simulator is closer to real biology. Vertical scales are per-panel. Improved-observable counts are reported per simulator against each one’s own baseline type (capability vs. deployment) and are never summed across the heterogeneous baselines; the cross-simulator-comparable statement is $G < 1$ (Table 2).

- [3] Cornelia I. Bargmann and Eve Marder. From the connectome to brain function. *Nature Methods*, 10(6):483–490, 2013.
- [4] Kristin Branson, Alice A. Robie, John Bender, Pietro Perona, and Michael H. Dickinson. High-throughput ethomics in large groups of *Drosophila*. *Nature Methods*, 6(6):451–457, 2009.
- [5] André E. X. Brown, Eviatar I. Yemini, Laura J. Grundy, Tadas Jucikas, and William R. Schafer. A dictionary of behavioral motifs reveals clusters of genes affecting *Caenorhabditis elegans* locomotion. *Proceedings of the National Academy of Sciences*, 110(2):791–796, 2013.
- [6] Vittorio Caggiano, Huawei Wang, Guillaume Durandau, Massimo Sartori, and Vikash Kumar. MyoSuite: a contact-rich simulation suite for musculoskeletal motor control. *Proceedings of Machine Learning Research*, 168:492–507, 2022.
- [7] Mohan Chen, Dazheng Feng, Hongtao Su, Tingwei Su, and Meng Wang. Neural model generating klinotaxis behavior accompanied by a random walk based on *C. elegans* connectome. *Scientific Reports*, 12(1):3043, 2022.
- [8] Steven J Cook, Travis A Jarrell, Christopher A Brittin, et al. Whole-animal connectomes of both *Caenorhabditis elegans* sexes. *Nature*, 571(7763):63–71, 2019.
- [9] Kyle Cranmer, Johann Brehmer, and Gilles Louppe. The frontier of simulation-based inference. *Proceedings of the National Academy of Sciences*, 117(48):30055–30062, 2020.
- [10] Marco Cuturi. Sinkhorn distances: Lightspeed computation of optimal transport. *Advances in Neural Information Processing Systems (NeurIPS)*, 26:2292–2300, 2013.
- [11] Scott L. Delp, Frank C. Anderson, Allison S. Arnold, Peter Loan, Ayman Habib, Chand T. John, Eran Guendelman, and Darryl G. Thelen. OpenSim: open-source software to create and analyze dynamic simulations of movement. *IEEE Transactions on Biomedical Engineering*, 54(11):1940–1950, 2007.
- [12] Michael H. Dickinson, Claire T. Farley, Robert J. Full, M. A. R. Koehl, Rodger Kram, and Steven Lehman. How animals move: an integrative view. *Science*, 288(5463):100–106, 2000.




Figure 3: **The PGOB formulation with an amortized neural-posterior engine (C4).** **a**, Posterior-predictive reproduction of each simulator’s own behaviour (modWorm mean absolute observable difference 1.4×10^{-3} ; FlyGym posterior-mean $\text{rel}_{\theta}=0.018$). **b**, Held-out negative log-probability for three density estimators (neural spline flow > MAF > mixture-density network on both simulators). The optimiser and the amortized posterior are interchangeable back-ends of one behaviour-only formulation.

- [13] Sven Dorkenwald, Arie Matsliah, Amy R. Sterling, et al. Neuronal wiring diagram of an adult brain. *Nature*, 634:124–138, 2024.
- [14] Conor Durkan, Artur Bekasov, Iain Murray, and George Papamakarios. Neural spline flows. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [15] Anthony D. Fouad, Shelly Teng, Julian R. Mark, Alice Liu, Pilar Alvarez-Illera, Hongfei Ji, et al. Distributed rhythm generators underlie *Caenorhabditis elegans* forward locomotion. *eLife*, 7:e29913, 2018.
- [16] Shangbang Gao, Sihui Asuka Guan, Anthony D. Fouad, Jun Meng, Taizo Kawano, Yung-Chi Huang, et al. Excitatory motor neurons are local oscillators for backward locomotion. *eLife*, 7:e29915, 2018.
- [17] Pdraig Gleeson, David Lung, Radu Grober, et al. c302: a multiscale framework for modelling the nervous system of *Caenorhabditis elegans*. *Philosophical Transactions of the Royal Society B*, 373(1758):20170379, 2018.
- [18] Pedro J Gonçalves, Jan-Matthis Lueckmann, Michael Deistler, et al. Training deep neural density estimators to identify mechanistic models of neural dynamics. *eLife*, 9:e56261, 2020.
- [19] Jesse M. Gray, Joseph J. Hill, and Cornelia I. Bargmann. A circuit for navigation in *caenorhabditis elegans*. *PNAS*, 102:3184–3191, 2005.
- [20] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, pages 1861–1870, 2018.
- [21] Nikolaus Hansen. The CMA evolution strategy: A tutorial. *arXiv preprint arXiv:1604.00772*, 2016.




Figure 4: **Behaviour-level generation of the *C. elegans* undulation wave (modWorm).** Body-bend-oscillation kymographs (temporal mean removed) for a random scrambled start, the PGOB-generated parameters, and the simulator’s ground-truth target ($n=8$ seeds each); the generated worm reproduces the travelling wave (bend-wave correlation to target $0.76 \rightarrow 1.00$), the random start does not. The match is behavioural: the generated parameters lie on the behaviourally-equivalent set, not necessarily at the exact ground-truth coordinates.

- [22] Nikolaus Hansen and Andreas Ostermeier. Completely derandomized self-adaptation in evolution strategies. *Evolutionary Computation*, 9(2):159–195, 2001.
- [23] Michael L. Hines and Nicholas T. Carnevale. The NEURON simulation environment. *Neural Computation*, 9(6):1179–1209, 1997.
- [24] Alan L. Hodgkin and Andrew F. Huxley. A quantitative description of membrane current and its application to conduction and excitation in nerve. *The Journal of Physiology*, 117(4): 500–544, 1952.
- [25] Auke Jan Ijspeert, Alessandro Crespi, Dimitri Ryczko, and Jean-Marie Cabelguen. From swimming to walking with a salamander robot driven by a spinal cord model. *Science*, 315 (5817):1416–1420, 2007.
- [26] Eduardo J. Izquierdo and Randall D. Beer. Connecting a connectome to behavior: an ensemble of neuroanatomical models of *C. elegans* klinotaxis. *PLoS Computational Biology*, 9(2):e1002890, 2013.
- [27] Eduardo J Izquierdo and Randall D Beer. Connecting a connectome to behavior: an ensemble of neuroanatomical models of *C. elegans* klinotaxis. *PLoS Computational Biology*, 14(2):e1005968, 2018.
- [28] Avelino Javier, Michael Currie, Chee Wai Lee, Jim Hokanson, Kezhi Li, Céline N. Martineau, Eviatar Yemini, Laura J. Grundy, Chris Li, QueeLim Ch’ng, William R. Schafer, Ellen A. A. Nollen, Rex Kerr, and André E. X. Brown. An open-source platform for analyzing and sharing worm-behavior data. *Nature Methods*, 15(9):645–646, 2018.




Figure 5: Scope of PGOB: parameter-space corruption is generated back, dynamics-space corruption is not. **a**, modWorm (Cook ODE) *dynamics*-space corruption: injecting state noise collapses observable closeness from 0.9999 to 0.30–0.46 (orange; grey are intact). **b**, FlyGym v2 (MuJoCo) *dynamics*-space corruption: the integrator diverges on 4/6 levels ($L_{final}=1.0$), leaving no usable forward model. **c**, BAAIWorm *parameter*-space corruption: corrupted-start loss L_0 vs. after-PGOB loss L_{final} per level; 30/30 cells improve, mean $\Delta L/L_0 = -29.4\%$ over the full L_0 – L_5 sweep. PGOB is a parameter-space generation framework, not a dynamics-space repair tool.

- [29] Pierre Karashchuk, Katie I. Rupp, Evyn S. Dickinson, Sarah Walling-Bell, Elischa Sanders, Eiman Azim, Bingni W. Brunton, and John C. Tuthill. Anipose: a toolkit for robust markerless 3D pose estimation. *Cell Reports*, 36(13):109730, 2021.
- [30] Jan Karbowski. Deciphering neural circuits for *Caenorhabditis elegans* behavior by computations and perturbations to genome and connectome. *Current Opinion in Systems Biology*, 13:44–51, 2019.
- [31] Jonathan R. Karr, Jayodita C. Sanghvi, Derek N. Macklin, Miriam V. Gutschow, Jared M. Jacobs, Benjamin Bolival Jr, Nacyra Assad-Garcia, John I. Glass, and Markus W. Covert. A whole-cell computational model predicts phenotype from genotype. *Cell*, 150(2):389–401, 2012.
- [32] Saul Kato, Harris S Kaplan, Tina Schrödel, et al. Global brain dynamics embed the motor command sequence of *Caenorhabditis elegans*. *Cell*, 163(3):656–669, 2015.
- [33] Janne K. Lappalainen, Fabian D. Tschopp, Sridhama Prakhya, Mason McGill, Aljoscha Nern, Kazunori Shinomiya, Shin-ya Takemura, Eyal Gruntman, Jakob H. Macke, and Srinivas C. Turaga. Connectome-constrained networks predict neural activity across the fly visual system. *Nature*, 634:1132–1140, 2024.
- [34] Victor Lobato-Rios, Shravan Tata Ramalingasetty, Pembe Gizem Özdil, Jonathan Arreguit, Auke Jan Ijspeert, and Pavan Ramdya. NeuroMechFly, a neuromechanical model of adult *Drosophila melanogaster*. *Nature Methods*, 19(5):620–627, 2022.
- [35] Johan M. Melis, Igor Siwanowicz, and Michael H. Dickinson. Machine learning reveals the control mechanics of an insect wing hinge. *Nature*, 628(8009):795–803, 2024.
- [36] Josh Merel, Diego Aldarondo, Jesse Marshall, Yuval Tassa, Greg Wayne, and Bence Ölveczky. Deep neuroethology of a virtual rodent. *arXiv preprint arXiv:1911.09451*, 2019. International Conference on Learning Representations (ICLR) 2020.

- [37] Florian T. Muijres, Michael J. Elzinga, Johan M. Melis, and Michael H. Dickinson. Flies evade looming targets by executing rapid visually directed banked turns. *Science*, 344(6180):172–177, 2014.
- [38] Jeffrey P. Nguyen, Frederick B. Shipley, Ashley N. Linder, George S. Plummer, Mochi Liu, Sagar U. Setru, Joshua W. Shaevitz, and Andrew M. Leifer. Whole-brain calcium imaging with cellular resolution in freely behaving *Caenorhabditis elegans*. *Proceedings of the National Academy of Sciences*, 113(8):E1074–E1081, 2016.
- [39] George Papamakarios and Iain Murray. Fast ε -free inference of simulation models with Bayesian conditional density estimation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2016.
- [40] George Papamakarios, Theo Pavlakou, and Iain Murray. Masked autoregressive flow for density estimation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [41] Xue Bin Peng, Pieter Abbeel, Sergey Levine, and Michiel van de Panne. DeepMimic: example-guided deep reinforcement learning of physics-based character skills. *ACM Transactions on Graphics*, 37(4):143, 2018.
- [42] Talmo D. Pereira, Joshua W. Shaevitz, and Mala Murthy. Quantifying behavior to understand the brain. *Nature Neuroscience*, 23(12):1537–1549, 2020.
- [43] Talmo D. Pereira, Nathaniel Tabris, Arie Matsliah, David M. Turner, Junyu Li, Shruthi Ravindranath, et al. SLEAP: a deep learning system for multi-animal pose tracking. *Nature Methods*, 19(4):486–495, 2022.
- [44] Gabriel Peyré and Marco Cuturi. Computational optimal transport: With applications to data science. *Foundations and Trends in Machine Learning*, 11(5–6):355–607, 2019.
- [45] Francesco Randi and Andrew M Leifer. Multi-modal whole-organism neuroscience: Integrating behavior, neural dynamics, and physiology in *C. elegans*. *Current Opinion in Neurobiology*, 79: 102684, 2023.
- [46] Andrea Saltelli, Marco Ratto, Terry Andres, Francesca Campolongo, Jessica Cariboni, Debora Gatelli, Michaela Saisana, and Stefano Tarantola. *Global Sensitivity Analysis: The Primer*. John Wiley & Sons, 2008.
- [47] Sanjay P. Sane. The aerodynamics of insect flight. *Journal of Experimental Biology*, 206(23): 4191–4208, 2003.
- [48] Gopal P. Sarma, Chee Wai Lee, Tom Portegys, Vahid Ghayoomie, Travis Jacobs, Bradley Alicea, et al. OpenWorm: overview and recent advances in integrative biological simulation of *Caenorhabditis elegans*. *Philosophical Transactions of the Royal Society B*, 373(1758):20170382, 2018.
- [49] Louis K. Scheffer, C. Shan Xu, Michal Januszewski, et al. A connectome and analysis of the adult *Drosophila* central brain. *eLife*, 9:e57443, 2020.
- [50] Eli Shlizerman et al. modworm: a modular muscular-skeletal *c. elegans* simulator. *Preprint*, 2025.

- [51] James C Spall. Multivariate stochastic approximation using a simultaneous perturbation gradient approximation. *IEEE Transactions on Automatic Control*, 37(3):332–341, 1992.
- [52] James C. Spall. *Introduction to Stochastic Search and Optimization: Estimation, Simulation, and Control*. John Wiley & Sons, Hoboken, NJ, 2003.
- [53] Greg J Stephens, Bethany Johnson-Kerner, William Bialek, and William S Ryu. Dimensionality and dynamics in the behavior of *C. elegans*. *PLoS Computational Biology*, 4(4):e1000028, 2008.
- [54] Greg J. Stephens, Matthew Bueno de Mesquita, William S. Ryu, and William Bialek. Emergence of long timescales and stereotyped behaviors in *Caenorhabditis elegans*. *Proceedings of the National Academy of Sciences*, 108(18):7286–7289, 2011.
- [55] Balázs Szigeti, Pádraig Gleeson, Michael Vella, et al. OpenWorm: An open-science approach to modeling *Caenorhabditis elegans*. *Frontiers in Computational Neuroscience*, 8:137, 2014.
- [56] Yuval Tassa, Yotam Doron, Alistair Muldal, Tom Erez, Yazhe Li, Diego de Las Casas, David Budden, Abbas Abdolmaleki, Josh Merel, Andrew Lefrancq, Timothy Lillicrap, and Martin Riedmiller. DeepMind control suite. *arXiv preprint arXiv:1801.00690*, 2018.
- [57] Alvaro Tejero-Cantero, Jan Boelts, Michael Deistler, Jan-Matthis Lueckmann, Conor Durkan, Pedro J Gonçalves, David S Greenberg, and Jakob H Macke. sbi: A toolkit for simulation-based inference. *Journal of Open Source Software*, 5(52):2505, 2020.
- [58] Emanuel Todorov, Tom Erez, and Yuval Tassa. MuJoCo: A physics engine for model-based control. In *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 5026–5033, 2012.
- [59] Saran Tunyasuvunakool, Alistair Muldal, Yotam Doron, Siqi Liu, Steven Bohez, Josh Merel, Tom Erez, Timothy Lillicrap, Nicolas Heess, and Yuval Tassa. dm_control: software and tasks for continuous control. *Software Impacts*, 6:100022, 2020.
- [60] Anne E. Urai, Brent Doiron, Andrew M. Leifer, and Anne K. Churchland. Large-scale neural recordings call for new insights to link brain and behavior. *Nature Neuroscience*, 25(1):11–19, 2022.
- [61] Lav R. Varshney, Beth L. Chen, Eric Paniagua, David H. Hall, and Dmitri B. Chklovskii. Structural properties of the *Caenorhabditis elegans* neuronal network. *PLoS Computational Biology*, 7(2):e1001066, 2011.
- [62] Roman Vaxenburg, Igor Siwanowicz, Josh Merel, Alice A Robie, Carmen Morrow, Guido Novati, Sam Stupski, Gwyneth M Card, and Srinivas C Turaga. Whole-body simulation of *drosophila* reveals motor control of behavior. *Nature*, 2025. flybody release, Figshare DOI 10.25378/janelia.25309105.
- [63] Alejandro F. Villaverde and Julio R. Banga. Reverse engineering and identification in systems biology: strategies, perspectives and challenges. *Journal of the Royal Society Interface*, 11(91):20130505, 2014.
- [64] Sibor Wang-Chen et al. Neuromechfly v2: simulating embodied sensorimotor control in adult *drosophila*. *Nature Methods*, 2024.

- [65] Quan Wen, Michelle D. Po, Elizabeth Hulme, Sway Chen, Xinyu Liu, Sze Wai Kwok, et al. Proprioceptive coupling within motor neurons drives *C. elegans* forward locomotion. *Neuron*, 76(4):750–761, 2012.
- [66] Quan Wen, Shangbang Gao, and Mei Zhen. *Caenorhabditis elegans* excitatory ventral cord motor neurons derive rhythm for body undulation. *Philosophical Transactions of the Royal Society B*, 373(1758):20170370, 2018.
- [67] John G White, Eileen Southgate, J Nichol Thomson, and Sydney Brenner. The structure of the nervous system of the nematode *Caenorhabditis elegans*. *Philosophical Transactions of the Royal Society of London B*, 314(1165):1–340, 1986.
- [68] Daniel Witvliet, Ben Mulcahy, James K. Mitchell, Yaron Meirovitch, Daniel R. Berger, Yuelong Wu, et al. Connectomes across development reveal principles of brain maturation. *Nature*, 596(7871):257–261, 2021.
- [69] Gang Yan, Petra E. Vértés, Emma K. Towlson, Yee Lian Chew, Denise S. Walker, William R. Schafer, and Albert-László Barabási. Network control principles predict neuron function in the *Caenorhabditis elegans* connectome. *Nature*, 550(7677):519–523, 2017.
- [70] Eviatar Yemini, Tadas Jucikas, Laura J. Grundy, André E. X. Brown, and William R. Schafer. A database of *Caenorhabditis elegans* behavioral phenotypes. *Nature Methods*, 10:877–879, 2013.
- [71] Anthony Zador, Sean Escola, Blake Richards, Bence Ölveczky, Yoshua Bengio, Kwabena Boahen, et al. Catalyzing next-generation artificial intelligence through NeuroAI. *Nature Communications*, 14:1597, 2023.
- [72] Mengdi Zhao, Ning Wang, Xinyu Jiang, Xiaohan Ma, Haoran Ma, Gan He, Kai Du, Lei Ma, and Tiejun Huang. An integrative data-driven model simulating *C. elegans* brain, body and environment interactions. *Nature Computational Science*, 4:106–120, 2024.